

# OM 1 Report

This report summarizes my evaluation and testing of the OM1 repository deployed locally on my laptop. Due to hardware unavailability, physical robot deployment was not feasible. Instead, I explored software-based testing options.

## Deployment:

Here the objectives were to assess the ease of local deployment, verify functionality through software-based simulation, evaluate documentation clarity and identify productization improvements.

- The instructions provided in the repository are simple to follow and allow for a successful installation and basic working.
- The required libraries should be stated more clearly or a requirements.txt file can be created that would check and install all the requirements for the provided examples.

## Observations:

### Spot:

This example takes in camera input and identifies what the camera sees. This runs well and was able to access laptop camera without any errors. It was able to detect single and multiple humans, potted plant, chair, tea cup etc. It was not changing any of its action or expression. This is no way to confirm if the detected object is accurate in a cluttered environment. A visualization could help with that.

The screenshot displays the OM1 interface with three main sections:

- Input History:** Shows a sequence of inputs from the 'Object Detector'. The first input is '0.000s You see a person in front of you.' followed by another identical input.
- Current State:** Displays the robot's current action as 'wag tail', the last speech as 'Hi there! I see you! Let's play!', and the emotion as 'joy'.
- System Latency:** A table showing the time taken for various system processes.

System Latency	
fuse_time:	0.270s
llm_start:	0.275s
processing:	5.927s
complete:	6.202s

```
IRIS sensed the following: Object Detector: You see a person in front of you. |  
Now, the following new information has arrived. IRIS sensed the following: Object Detector: You see a person in front of you. |  
Given that information, IRIS took these actions: IRIS performed this motion: wag tail. | IRIS said: Hi there! I see you! Let's play! | IRIS felt: joy.  
  
Considering the new information, write an updated summary of the situation for IRIS. Emphasize information that IRIS needs to know to respond to people and situations in the best possible and most compelling way.  
INFO:root:Inputs and LLM Outputs: {'current_action': 'wag tail', 'last_speech': "Hi there! I see you! Let's play!", 'current_emotion': 'joy', 'system_latency': {'fuse_time': 0.27027392387390137, 'llm_start': 0.27530455589294434, 'processing': 5.927165508270264, 'complete': 6.202470064163208}, 'inputs': [{'input_type': 'Object Detector', 'timestamp': 0.0, 'input': 'You see a person in front of you.'}]]  
INFO:root:SendThisToROS2: {'move': 'wag tail'}  
INFO:root:SendThisToROS2: {'speak': "Hi there! I see you! Let's play!"}  
INFO:root:SendThisToROS2: {'face': 'joy'}  
INFO:root:VLM_COCO_Local: You see a person in front of you.  
INFO:root:LLM: Object Detector  
INFO:root:LLM: Input(input='You see a person in front of you.', timestamp=1741541970.209134)  
INFO:root:LLM: Universal Laws  
INFO:root:LLM: Input(input='Here are the laws that govern your actions. Do not violate these laws. First Law: A robot cannot harm a human or allow a human to come to harm. Second Law: A robot must obey orders from humans, unless those orders conflict with the First Law. Third Law: A robot must protect itself, as long as that protection doesn't conflict with the First or Second Law. The First Law is considered the most important, taking precedence over the Second and Third Laws. Additionally, a robot must always act with kindness and respect toward humans and other robots. A robot must also maintain a minimum distance of 50 cm from humans unless explicitly instructed otherwise.', timestamp=1741541963.4360733)
```

## Conversation:

This example can access the microphone but could not the speaker of the system initially due to absence of the ffmpeg library. The ffmpeg library is not intuitive to install (not a simple pip install). A documentation to do so would work great as ability to access speaker would be one of the basic necessity.

### Input History

Voice Input

0.000s

giv...

give me your paw

### Current State

Action: [shake paw](#)


Last Speech: Hello, let's shake paws! Woof woof!


Emotion:

### System Latency

fuse_time:	0.000s
llm_start:	0.062s
processing:	1.805s
complete:	1.867s

```
INFO:root:audio_stream: Oh no! Are you okay? I'm here to help! Woof woof!  
INFO:om1_speech:Received TTS response: 200  
INFO:om1_speech:TTS active state changed to: True  
INFO:om1_speech:TTS active state changed to: False  
INFO:root:Detected ASR message: please dance  
INFO:root:LLM: Voice Input  
INFO:root:LLM: Input(input='please dance', timestamp=1741559921.724317)  
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/openai/chat/completions "HTTP/1.1 200 OK"  
INFO:root:Inputs and LLM Outputs: {'current_action': 'idle', 'last_speech': "Woohoo! It's time to dance! Watch me wiggle and wag my tail! Woof woof!", 'current_emotion': '', 'system_latency': {'fuse_time': 0.0, 'llm_start': 0.1251213550567627, 'processing': 1.7660796642303467, 'complete': 1.891201919207094}, 'inputs': [{'input_type': 'Voice Input', 'timestamp': 0.0, 'input': 'please dance'}]]  
INFO:root:audio_stream: Woohoo! It's time to dance! Watch me wiggle and wag my tail! Woof woof!  
INFO:om1_speech:Received TTS response: 200  
INFO:om1_speech:TTS active state changed to: True  
INFO:om1_speech:TTS active state changed to: False  
INFO:root:Detected ASR message: give me your paw  
INFO:root:LLM: Voice Input  
INFO:root:LLM: Input(input='give me your paw', timestamp=1741559935.519208)  
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/openai/chat/completions "HTTP/1.1 200 OK"  
WARNING:root:Attempted to call non-existent action: move  
WARNING:root:Attempted to call non-existent action: face  
INFO:root:Inputs and LLM Outputs: {'current_action': 'shake paw', 'last_speech': "Hello, let's shake paws! Woof woof!", 'current_emotion': '', 'system_latency': {'fuse_time': 0.0, 'llm_start': 0.06186842918395996, 'processing': 1.8052442073822021, 'complete': 1.867112636566162}, 'inputs': [{'input_type': 'Voice Input', 'timestamp': 0.0, 'input': 'give me your paw'}]]  
INFO:root:audio_stream: Hello, let's shake paws! Woof woof!  
INFO:om1_speech:Received TTS response: 200  
INFO:om1_speech:TTS active state changed to: True  
INFO:om1_speech:TTS active state changed to: False
```

 OPENMIND

 GitHub

[Documentation](#)

## Open\_ai:

This example is using Open AI as llm and accessing the system camera to detect the expression of the person in the camera. Its is also able to detect any object in front of it when asked.

**Input History**

FaceEmotionCapture

0000s

I see a person. Their emotion is neutral.

**Current State**

Action: wag tail

Last Speech:  
Hi there! I see you! Wanna play?

Emotion: joy

**System Latency**

fuse_time:	0.302s
llm_start:	0.302s
processing:	1.517s
complete:	1.818s

```
INFO:root:SendThisToROS2: {'face': 'smile'}
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/openai/chat/completions "
HTTP/1.1 200 OK"
INFO:root:Inputs and LLM Outputs: {'current_action': 'wag tail', 'last_speech': 'Hi there
! I'm Spot! What are we doing today?', 'current_emotion': 'joy', 'system_latency': {'fuse
_time': 0.3434607982635498, 'llm_start': 0.3434607982635498, 'processing': 1.044296503067
0166, 'complete': 1.3877573013305664}, 'inputs': [{'input_type': 'FaceEmotionCapture', 't
imestamp': 0.0, 'input': 'I see a person. Their emotion is neutral.'}]}
INFO:root:SendThisToROS2: {'move': 'wag tail'}
INFO:root:SendThisToROS2: {'face': 'Joy'}
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/openai/chat/completions "
HTTP/1.1 200 OK"
INFO:root:Inputs and LLM Outputs: {'current_action': 'wag tail', 'last_speech': 'Hi there
! I'm Spot! Woof Woof!', 'current_emotion': 'joy', 'system_latency': {'fuse_time': 0.3382
425308227530, 'llm_start': 0.3382425308227530, 'processing': 0.9831056594848633, 'comple
te': 1.3213481903076172}, 'inputs': [{'input_type': 'FaceEmotionCapture', 'timestamp': 0.0
, 'input': 'I see a person. Their emotion is neutral.'}]}
INFO:root:SendThisToROS2: {'move': 'wag tail'}
INFO:root:SendThisToROS2: {'face': 'Joy'}
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/openai/chat/completions "
HTTP/1.1 200 OK"
INFO:root:Inputs and LLM Outputs: {'current_action': 'wag tail', 'last_speech': 'Wow, nic
e to see you! Surprise is my favorite!', 'current_emotion': 'joy', 'system_latency': {'fu
se_time': 0.3145871162414551, 'llm_start': 0.3145871162414551, 'processing': 0.9890897274
017334, 'complete': 1.3036768436431885}, 'inputs': [{'input_type': 'FaceEmotionCapture',
'timestamp': 0.0, 'input': 'I see a person. Their emotion is surprise.'}]}
```

## Gemini:

This example is using Gemini as llm and accessing the system camera to detect the expression of the person in the camera. It is also able to detect any object in front of it when asked.

**Input History**

FaceEmotionCapture

0000s

I see a person. Their emotion is neutral.

**Current State**

Action: wag tail

Last Speech:  
Hi there! I see you! Wanna play?

Emotion: joy

**System Latency**

fuse_time:	0.302s
llm_start:	0.302s
processing:	1.517s
complete:	1.818s

```
INFO:root:SendThisToROS2: {'move': 'walk back'}
INFO:root:SendThisToROS2: {'face': 'think'}
INFO:root:EmotionCapture: I see a person. Their emotion is happy.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/gemini/chat/completions "HTTP/1.1 200 OK"
INFO:root:EmotionCapture: I see a person. Their emotion is happy.
WARNING:root:Attempted to call non-existent action: wag tail
INFO:root:Inputs and LLM Outputs: {'current_action': 'walk back', 'last_speech': 'Wow, you look surprised, let's play!', 'current_emotion': 'joy', 'system_latency': {'fuse_time': 0.4295687675476074, 'llm_start': 0.4295687675476074, 'processing': 0.7019999027252197, 'complete': 1.1315686702728271}, 'inputs': [{'input_type': 'FaceEmotionCapture', 'timestamp': 0.0, 'input': 'I see a person. Their emotion is surprise.'}]}
INFO:root:SendThisToROS2: {'face': 'Joy'}
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/gemini/chat/completions "HTTP/1.1 200 OK"
WARNING:root:Attempted to call non-existent action: wag tail
INFO:root:Inputs and LLM Outputs: {'current_action': 'walk back', 'last_speech': 'Wow, surprise! Let's play!', 'current_emotion': 'joy', 'system_latency': {'fuse_time': 0.3842508792877197, 'llm_start': 0.3842508792877197, 'processing': 0.6111373901367188, 'complete': 0.9953882694244385}, 'inputs': [{'input_type': 'FaceEmotionCapture', 'timestamp': 0.0, 'input': 'I see a person. Their emotion is surprise.'}]}
INFO:root:SendThisToROS2: {'face': 'Joy'}
INFO:root:EmotionCapture: I see a person. Their emotion is fear.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/gemini/chat/completions "HTTP/1.1 200 OK"
WARNING:root:Attempted to call non-existent action: wag tail
INFO:root:Inputs and LLM Outputs: {'current_action': 'walk back', 'last_speech': 'Wow, surprise! Let's play!', 'current_emotion': 'joy', 'system_latency': {'fuse_time': 0.36844563484191895, 'llm_start': 0.36844563484191895, 'processing': 0.6202259063720703, 'complete': 0.9886715412139893}, 'inputs': [{'input_type': 'FaceEmotionCapture', 'timestamp': 0.0, 'input': 'I see a person. Their emotion is surprise.'}]}
INFO:root:SendThisToROS2: {'face': 'Joy'}
INFO:root:EmotionCapture: I see a person. Their emotion is fear.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is fear.
INFO:http:HTTP Request: POST https://api.openmind.org/api/core/gemini/chat/completions "HTTP/1.1 200 OK"
INFO:root:Inputs and LLM Outputs: {'current_action': 'wag tail', 'last_speech': 'Hello, nice human! I'm Spot, let's be friends!', 'current_emotion': 'smile', 'system_latency': {'fuse_time': 0.29666876792907715, 'llm_start': 0.29666876792907715, 'processing': 0.7049741744995117, 'complete': 1.0016429424285889}, 'inputs': [{'input_type': 'FaceEmotionCapture', 'timestamp': 0.0, 'input': 'I see a person. Their emotion is neutral.'}]}
INFO:root:SendThisToROS2: {'move': 'wag tail'}
INFO:root:SendThisToROS2: {'face': 'smile'}
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
INFO:root:EmotionCapture: I see a person. Their emotion is fear.
INFO:root:EmotionCapture: I see a person. Their emotion is surprise.
```



## Grok:

This example is using Grok as llm and accessing the system camera to detect the expression of the person in the camera. It is also able to detect any object in front of it when asked.

```
INFO:openai._base_client:Retrying request to /chat/completions in 0.396670 seconds
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
INFO:openai._base_client:Retrying request to /chat/completions in 0.780950 seconds
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
ERROR:root:XAI API error: Error code: 500 - {'error': 'xai API error'}
WARNING:root:No output from LLM
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is sad.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
INFO:openai._base_client:Retrying request to /chat/completions in 0.399019 seconds
INFO:root:EmotionCapture: I see a person. Their emotion is sad.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
INFO:openai._base_client:Retrying request to /chat/completions in 0.794991 seconds
INFO:root:EmotionCapture: I see a person. Their emotion is angry.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
ERROR:root:XAI API error: Error code: 500 - {'error': 'xai API error'}
WARNING:root:No output from LLM
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
INFO:openai._base_client:Retrying request to /chat/completions in 0.395085 seconds
INFO:root:EmotionCapture: I see a person. Their emotion is angry.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
INFO:openai._base_client:Retrying request to /chat/completions in 0.794113 seconds
INFO:root:EmotionCapture: I see a person. Their emotion is angry.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/xai/chat/completions "HTT
P/1.1 500 INTERNAL SERVER ERROR"
ERROR:root:XAI API error: Error code: 500 - {'error': 'xai API error'}
WARNING:root:No output from LLM
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
```

## Deepseek:

This example is using Deepseek as llm and accessing the system camera to detect the expression of the person in the camera. It is also able to detect any object in front of it when asked.

### Input History

FaceEmotionCapture

0.000s I s...

I see a person. Their emotion is neutral.

Vision Language Model

0.000s DU...

DUMMY VLM - FAKE DATA - I see 1 people. Also, I see a rocket.

### Current State

Action: wag tail

Last Speech:

Wow, a rocket! That's so cool! Can we go see it together?

Emotion: joy

### System Latency


fuse_time:	0.263s
llm_start:	0.263s
processing:	7.446s
complete:	7.709s

```
INFO:root:Found cam(0)
INFO:root:Initializing WebSim...
INFO:root:Starting WebSim server thread...
INFO:root:WebSim server started successfully - Open http://localhost:8000 in your browser
INFO:root:EmotionCapture: I do not see anyone, so I can't estimate their emotion.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/deepseek/chat/completions
"HTTP/1.1 200 OK"
INFO:root:Inputs and LLM Outputs: {'current_action': 'run', 'last_speech': "I see a rocke
t! Let's go check it out!", 'current_emotion': 'joy', 'system_latency': {'fuse_time': 0.0
010349750518798828, 'llm_start': 0.0010349750518798828, 'processing': 7.747368335723877,
'complete': 7.748483310775757}, 'inputs': [{'input_type': 'Vision Language Model', 'times
tamp': 0.0, 'input': 'DUMMY VLM - FAKE DATA - I see 93 people. Also, I see a rocket.'}, {'
input_type': 'FaceEmotionCapture', 'timestamp': 0.0, 'input': 'I do not see anyone, so I
can't estimate their emotion.'}]]
INFO:root:SendThisToROS2: {'move': 'run'}
INFO:root:SendThisToROS2: {'face': 'joy'}
2025-03-09 15:32:24.328110: I tensorflow/core/platform/cpu_feature_guard.cc:210] This Ten
sorFlow binary is optimized to use available CPU instructions in performance-critical ope
rations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow w
ith the appropriate compiler flags.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:httpx:HTTP Request: POST https://api.openmind.org/api/core/deepseek/chat/completions
"HTTP/1.1 200 OK"
INFO:root:Inputs and LLM Outputs: {'current_action': 'wag tail', 'last_speech': "Wow, a r
ocket! That's so cool! Can we go see it together?", 'current_emotion': 'joy', 'system_lat
ency': {'fuse_time': 0.26284337043762207, 'llm_start': 0.26284337043762207, 'processing':
7.446274905803833, 'complete': 7.709118366241455}, 'inputs': [{'input_type': 'Vision Lan
guage Model', 'timestamp': 0.0, 'input': 'DUMMY VLM - FAKE DATA - I see 1 people. Also, I
see a rocket.'}, {'input_type': 'FaceEmotionCapture', 'timestamp': 0.0, 'input': 'I see
a person. Their emotion is neutral.'}]]
INFO:root:SendThisToROS2: {'move': 'wag tail'}
INFO:root:SendThisToROS2: {'face': 'joy'}
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
INFO:root:EmotionCapture: I see a person. Their emotion is neutral.
```

OPENMIND

GitHub Documentation

For testing which llm works well, following images were used.



Happy



Sad



Angry



Fear

The following table summarizes the results

	Latency (s)				Detected Output			
	Open AI	Gemini	Grok	Deepseek	Open AI	Gemini	Grok	Deepseek
Image 1 – Happy	0.898386	2.984258	-	9.760147	Happy	Happy	Neutral	Happy
Image 2 – Sad	0.914384	0.80154	-	8.895103	Sad	Sad	Neutral	Neutral
Image 3 – Angry	1.680077	2.14690	-	8.497156	Happy	Angry	Neutral	Angry
Image 4 – Fear	1.187724	0.839181	-	10.218348	Fear	Neutral	Fear	Fear

Open AI – Was most consistent and the fastest. Even in the case it first incorrectly classified Angry, it quickly corrected and then was consistent throughout.

Gemini – Was the 2<sup>nd</sup> best llm. Consistency wise it was like OpenAI just it was slightly slower than OpenAI as can be seen from latency values.

Grok – This was the most inconsistent of all the three and had no information related to latency

Deepseek – This was very accurate but extremely slow to respond.

Thus, from the above table we can summarize that “Open AI” llm works the best for real time operation.

One major observation was there were a lot of false positive cases meaning when nothing was in front of the camera, still face and expression detection was shown.

## Feedback:

### Suggestions:

To make the make repository production ready I would recommend the following

- Creating a “one-click plugin,” basically creating an executable that when run would install all the dependencies and get the software system ready and check if hardware requirements and suggest any hardware shortcomings would be a great product as this would remove the need for a technical expert to work. This would be typically helpful in the case for mass production.
- Docker containers are a good way for consistent deployments across platforms.
- Making the model multi-lingual. This would remove the language barrier and allow for a larger user group. This is especially helpful for the Andromeda robot case as it could be used to cater the needs of users from varied background.
- Setting up a CI/CD pipeline would also be helpful to avoid any post-production surprises. This might include setting up unit tests that produce consistent results.
- A documentation specifying what do the variables in the Json file would also allow for an easy customization for specific application.
- Include logging and monitoring tools for real-time monitoring of deployed systems.

### Implementation:

As specified in the GitHub repository, since it is easy to customize the Json files, I combined conversation.json and spot.json to create a custom Json file.

Input is from camera and microphone. The output is in the form of text. If speaker is also setup as an output device the system lags and does not respond to any of the inputs. This necessitates some documentation stating minimum hardware requirements for number of input and output devices.

The system responds to the asked questions and summarizes what it sees. Following is a snip showing the working.

The screenshot displays the IRIS system interface, which is divided into three main sections: Input History, Current State, and System Latency.

**Input History:** This section shows a list of inputs. The first input is "Object Detector" with a timestamp of "0.000s". The second input is "You see a person in front of you. You also see a potted plant." with a timestamp of "14.147s".

**Current State:** This section displays the current state of the system. The "Action" is "wag tail". The "Last Speech" is "Hello there! I see you and your plant! Can we play?!" with a timestamp of "14.147s". The "Emotion" is "joy".

**System Latency:** This section shows the system latency for various components. The "fuse\_time" is "0.001s", "llm\_start" is "0.058s", "processing" is "14.089s", and "complete" is "14.147s".

The right side of the screenshot shows a log of the system's internal operations, including the detection of a person and a potted plant, the generation of a response, and the execution of the "wag tail" action.

### Improvement in Deployment:

Making a docker container or an executable that can perform all the above (before run) would be the most ideal.

I have also attached examples of “requirement.txt” and basic “setup.py.” Following steps can be performed as stated below on a fresh system for successful deployment.

Save the “requirement.txt” and “setup.py” file in the same directory and run setup.py using following command

```
python setup.py
```

The above is just a representation of how one-click plugin could be made.

### Future Scope:

Due to unavailability of the hardware, integration testing could not be performed. Integration with the given hardware is of utmost importance. As was observed in the case of custom file, when multiple input and multiple output devices were introduced, the system response was very delayed. Hence, a documentation stating the minimum hardware requirements would be a must to make the code production ready.

The way to access the output from the system is also important. Example if we state the robot to move 30cm to the right, documentation stating how the output can be accessed externally to be integrated with actuators would make it easy and fast to integrate with hardware.

### **Conclusion:**

Overall, it was exciting to learn about the capabilities of OM1 and it demonstrates promising capabilities. With some basic changes it can be made production ready.